

# Resource Allocation in Cloud and Fog Computing

Fog and Cloud Computing

---

**Francescomaria Faticanti**, Inria Sophia Antipolis, France

May 24, 2022

`francescomaria.faticanti@inria.fr`

# Why Resource Allocation?

Several benefits

- Save money
- Improve productivity
- Saving time
- Effective use of resources

# Why Resource Allocation?

Several benefits

- Save money
- Improve productivity
- Saving time
- Effective use of resources

How hard is that?

# Resource Allocation Problem

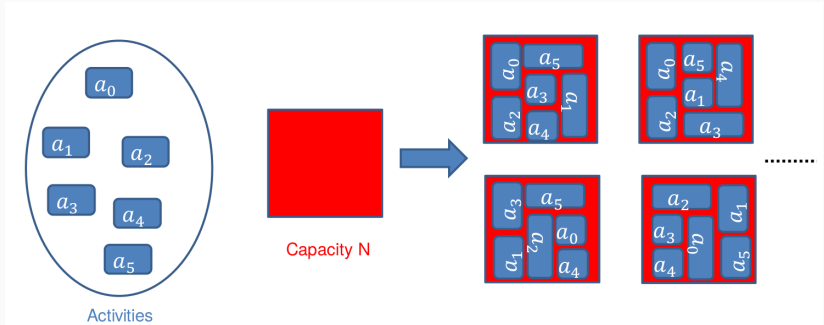
Given a type of resource whose total amount is equal to  $N$ , we want to allocate it to  $n$  activities so that a certain **objective function** is minimised (or maximised)

# Resource Allocation Problem

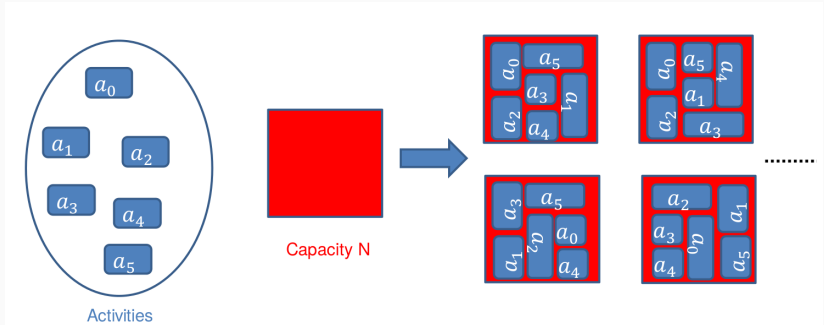
Given a type of resource whose total amount is equal to  $N$ , we want to allocate it to  $n$  activities so that a certain **objective function** is minimised (or maximised)

- **Objective function:** cost or loss (profit or reward) of the resulting allocation

# Resource Allocation Problem (1)



# Resource Allocation Problem (1)



Each configuration: specific objective value

*minimize*  $f(x_1, x_2, \dots, x_n)$

*s. t.*

$$(1) \sum_{\{j=1\}}^n x_j = N,$$

$$(2) \quad x_j \geq 0, \quad j = 1, 2, \dots, n.$$

**Constraint 1:**  
the total amount  
reserved to all  
the activities  
must be equal  
to the available  
resource

**Objective  
function:** its  
value depends on  
the values  
assigned to the  
variables

**Constraint 2:** the amount  
reserved to each activity  
must be nonnegative



## Distribution of search effort:

- Detect the position of an object whose possible positions are among those numbered from 1 to  $n$
- The probability of the object being in position  $j$  is  $p_j$
- $(1 - e^{-\alpha x_j}) p_j$ : conditional probability of detecting the object in position  $j$
- $x_j$ : amount of search effort
- **Constraint:**  $N$  is the total amount of effort
- **Objective:** maximize the overall probability of detecting the object

## Resource Distribution Problem:

- $n$  locations to which resources such as newspapers are distributed from the central factory
- Demand at each location  $j$
- $q_j(x_j)$ : the expected cost at location  $j$  when  $x_j$  is allocated to  $j$
- $N$ : total amount of resources allocated to all locations

$$\begin{aligned} & \text{minimize} \quad \sum_{j=1}^n q_j(x_j) \\ & \text{subject to} \quad \sum_{j=1}^n x_j = N, \\ & \quad \quad \quad x_j \geq 0. \end{aligned}$$

## Online Caching Problem

- Requests from a catalog of  $N$  files  $\mathcal{N} = \{1, \dots, N\}$
- $w_i \in \mathbb{R}_+$ : cost to serve a request for file  $i$  by a remote server
- Single-cache system with capacity  $K$
- Storage: arbitrary fractions of files
- Slotted time  $t \in \mathcal{T} = \{1, \dots, T\}$
- Cache state  $\mathbf{x}_t = [x_{t,i}]_{i \in \mathcal{N}}$  drawn from

$$\mathcal{X} = \left\{ \mathbf{x} \in [0, 1]^N : \sum_{i=1}^N x_i = K \right\},$$

where  $x_{i,t}$ : fraction of file  $i$  stored at time  $t$

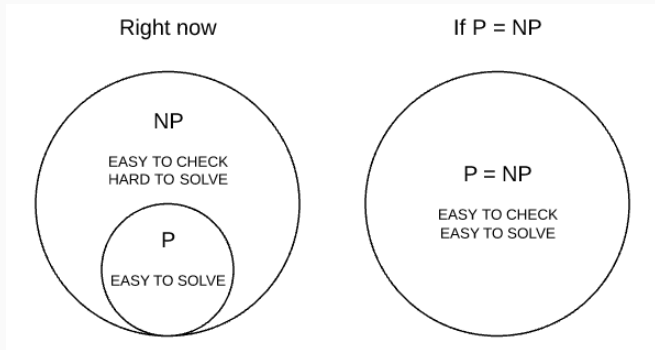
## Online Caching Problem (1)

$$\min_{\mathbf{x} \in \mathcal{X}} \sum_{t=1}^T f_t(\mathbf{x}) \quad (1)$$

$$\text{s.t. } \sum_{i=1}^N x_{t,i} = K, \quad \forall t \in \mathcal{T} \quad (2)$$

$$0 \leq x_{t,i} \leq 1, \quad \forall i \in \mathcal{N}, \forall t \in \mathcal{T} \quad (3)$$

# P vs NP



**Millennium Prize Problem**

# Decision Problem

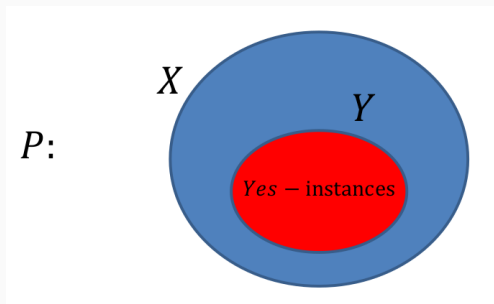
## Definition (Decision Problem)

A **decision** problem is a pair  $P = (X, Y)$ , where  $X$  is a language decidable in polynomial time and  $Y \subseteq X$ . Elements of  $X$ : **instances** of  $P$ ; elements of  $Y$ : **yes-instances**.

# Decision Problem

## Definition (Decision Problem)

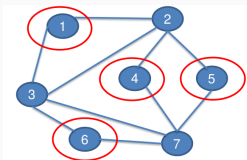
A **decision** problem is a pair  $P = (X, Y)$ , where  $X$  is a language decidable in polynomial time and  $Y \subseteq X$ . Elements of  $X$ : **instances** of  $P$ ; elements of  $Y$ : **yes-instances**.



## Example (Decision Problem)

### Problem (Independent Set)

*Given a graph  $G$ , a  $k$ -independent set is a subset of its vertices of cardinality  $k$ , where no couple of vertices is connected by an edge in  $G$ .*



4-independent set

- Instances:  $\langle G, k \rangle$
- Yes-instances:  $\langle G, k \rangle$  s.t.  $G$  contains an independent set of  $k$  vertices

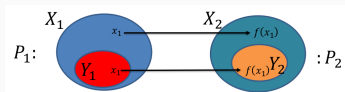


# Polynomial Reduction

## Definition (Polynomial Transformation)

Let  $P_1 = (X_1, Y_1)$  and  $P_2 = (X_2, Y_2)$  be decision problems.  $P_1$  **polynomially transforms** to  $P_2$  ( $P_1 \leq_p P_2$ ) if there is a function  $f : X_1 \rightarrow X_2$ , computable in polynomial time, s.t.

$$f(x_1) \in Y_2, \forall x_1 \in Y_1 \wedge f(x_1) \in X_2 \setminus Y_2, \forall x_1 \in X_1 \setminus Y_1.$$



## Proposition

*If  $P_1 \leq_p P_2$  and there is a polynomial time algorithm for  $P_2$ , then there is a polynomial time algorithm for  $P_1$ .*

## Definition

A problem  $P$  is **NP-complete** if

- $P \in NP$
- $\forall P' \in NP, P' \leq_p P$  (NP-hardness)

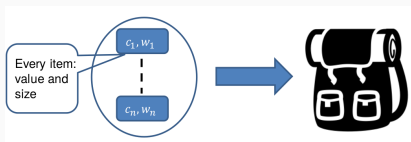
## Theorem

*SATISFIABILITY is NP-complete (Cook [1971])*

# NP-hard Problems

## Knapsack:

- *Instance:* A number  $n \in \mathbb{N}$  and nonnegative integers  $c_i$ ,  $w_i$  and  $W$  for  $i = 1, \dots, n$ .
- *Task:* Find a subset  $S \subseteq \{1, \dots, n\}$  such that  $\sum_{i \in S} w_i \leq W$  and  $\sum_{i \in S} c_i$  is maximum.

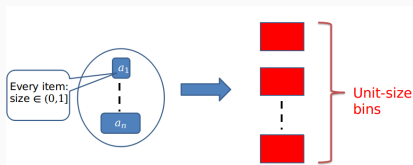


- Maximize the total value of collected items without exceeding the knapsack's capacity

# NP-hard Problems

## Bin-Packing:

- *Instance*: A list of nonnegative numbers  $a_1, \dots, a_n \leq 1$ .
- *Task*: Find a  $k \in \mathbf{N}$  and an assignment  $f : \{1, \dots, n\} \rightarrow \{1, \dots, k\}$  with  $\sum_{i:f(i)=j} a_i \leq 1, \forall j \in \{1, \dots, k\}$  s.t.  $k$  is minimum.

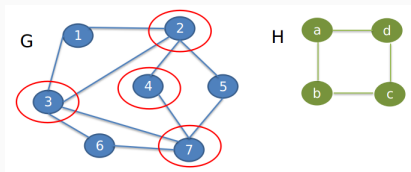


- Packing all the items using the minimum number of bins

# NP-hard Problems

## Subgraph Isomorphism:

- *Instance:* Two graphs  $G = (V, E)$  and  $H = (V', E')$ .
- *Task:* Find a subgraph  $G_0 = (V_0, E_0) : V_0 \subseteq V, E_0 \subseteq E$  s.t.  $G_0 \cong H$ .



- Does G contains a subgraph isomorphic to H?

# Challenge

- **Theory:** We cannot find poly-time algorithms for NP-hard problems (unless  $P = NP$ )
- **Desired features:**
  - Optimality
  - Poly-time
  - Arbitrary instances of the problem
- **Trade-off:**
  - Quality of the solution
  - Efficiency

## Approximation algorithms

## Approximation algorithms

- Approximation guarantees



## Approximation algorithms

- Approximation guarantees
- Approximation proof

## Approximation algorithms

- Approximation guarantees
- Approximation proof
- General case

## Approximation algorithms

- Approximation guarantees
- Approximation proof
- General case
- Efficient

## Heuristics

## Approximation algorithms

- Approximation guarantees
- Approximation proof
- General case
- Efficient

## Heuristics

- No approximation guarantees

## Approximation algorithms

- Approximation guarantees
- Approximation proof
- General case
- Efficient

## Heuristics

- No approximation guarantees
- Experimental proof of approximation

## Approximation algorithms

- Approximation guarantees
- Approximation proof
- General case
- Efficient

## Heuristics

- No approximation guarantees
- Experimental proof of approximation
- Specific case (most of the time)

## Approximation algorithms

- Approximation guarantees
- Approximation proof
- General case
- Efficient

## Heuristics

- No approximation guarantees
- Experimental proof of approximation
- Specific case (most of the time)

**Note:** hard to find good approximation algorithms for several NP-hard problems

**Online Bin-Packing** Consider the items in a given order and place them one by one.

Several heuristics:

- **First-Fit**
- **Best-Fit**
- **Worst-Fit**
- **Next-Fit**

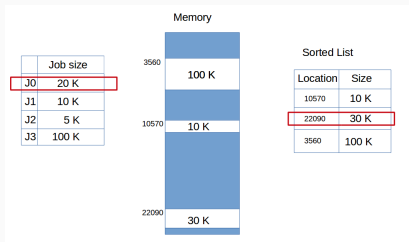


## Memory Allocation (Best-Fit)

- *Input*: A fixed-size job and the list of free-memory holes.
- *Algorithm*: It keeps the free-memory list in increasing order by size. The OS allocates the input job to the first fitting free partition in the sorted list.

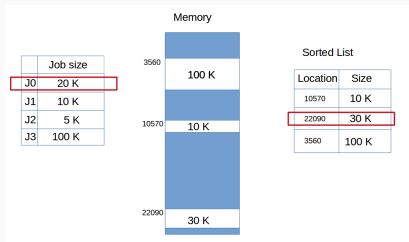
# Memory Allocation (Best-Fit)

- *Input:* A fixed-size job and the list of free-memory holes.
- *Algorithm:* It keeps the free-memory list in increasing order by size. The OS allocates the input job to the first fitting free partition in the sorted list.



# Memory Allocation (Best-Fit)

- *Input:* A fixed-size job and the list of free-memory holes.
- *Algorithm:* It keeps the free-memory list in increasing order by size. The OS allocates the input job to the first fitting free partition in the sorted list.



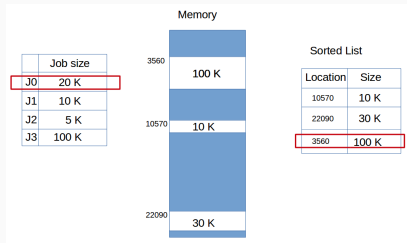
Pros and cons?

## Memory Allocation (Worst-Fit)

- *Input*: A fixed-size job and the list of free-memory holes.
- *Algorithm*: The OS allocates the input job to the first fitting free partition with the maximum size in the list.

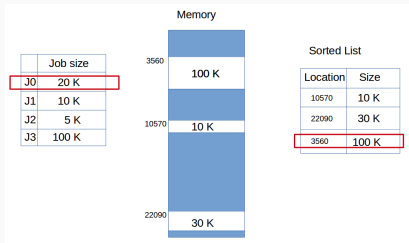
# Memory Allocation (Worst-Fit)

- *Input:* A fixed-size job and the list of free-memory holes.
- *Algorithm:* The OS allocates the input job to the first fitting free partition with the maximum size in the list.



# Memory Allocation (Worst-Fit)

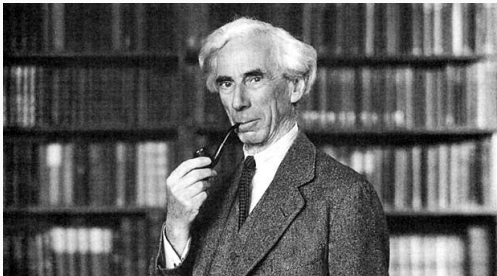
- *Input:* A fixed-size job and the list of free-memory holes.
- *Algorithm:* The OS allocates the input job to the first fitting free partition with the maximum size in the list.



Pros and cons?

# Approximation

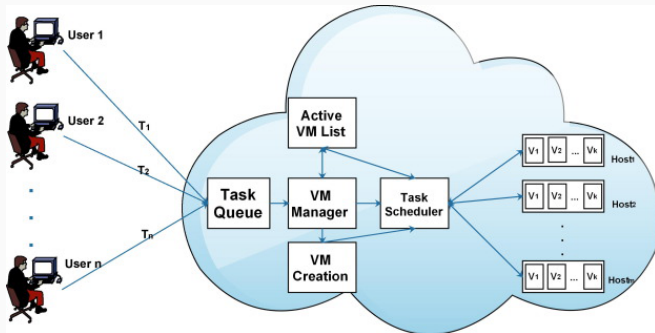
*"Although this may seem a paradox, all exact science is dominated by the idea of approximation."* **Bertrand Russell**



# Resource allocation in the Cloud



# Cloud



# Resource Allocation in Cloud

- Applications run on VMs
- Virtualization: splitting and assigning physical servers to VMs
- Active servers: main energy consumer in data centers
- Energy consumption: up to 50% or more of the total operational costs
- **Problem:** assigning VMs with volatile demands to physical servers such that energy costs are minimized
- **Observation:** minimizing energy costs → minimizing the number of active servers

## Typical approaches:

- Static Allocation
  - Optimization
  - Approximation
- Dynamic Allocation
  - Optimal Control techniques
  - Service live migration

- **Assumption:** workloads patterns are known
- **Challenge:** NP-hardness
- **Two main approaches:**
  - Mathematical Programming
  - Approximated Solutions

- **Pros:**
  - Mathematical formalization
  - Efficient solvers for problem resolution
- **Cons:**
  - Not scalable (as the instance size grows)
  - Most of the problems are NP-hard

# Examples

- **Bin-packing problem:** Given  $n$  items with sizes  $a_1, \dots, a_n \in (0, 1]$ , find a packing in unit-sized bins that minimizes the number of used bins.

$$\begin{aligned} & \text{minimize} \quad \sum_{j=1}^n y_j \\ & \text{s. t.} \quad \sum_{i=1}^n a_i x_{ij} \leq y_j, \quad j = 1, \dots, n; \\ & \quad \sum_{j=1}^n x_{ij} = 1, \quad i = 1, \dots, n; \\ & \quad y_i, x_{ij} \in \{0, 1\}, \quad i, j = 1, \dots, n. \end{aligned}$$

Each item must be packed in exactly one bin

Equal to 1 if bin  $j$  is used

The capacity of each bin cannot be exceeded

- **Server Consolidation problem:** The IT service manager needs to consolidate servers by assigning  $n$  services to target  $m$  servers minimizing the number of servers used.

# Examples

- **Server Consolidation problem:** The IT service manager needs to consolidate servers by assigning  $n$  services to target  $m$  servers minimizing the number of servers used.

$$\begin{aligned} & \text{minimize } \sum_{i=1}^m c_i y_i \\ \text{s.t. } & \sum_{j=1}^n u_{jk} x_{ij} \leq s_{ik} y_i, \quad \forall i \in I, \forall k \in K; \\ & \sum_{i=1}^m x_{ij} = 1, \quad \forall j \in J; \\ & y_i, x_{ij} \in \{0,1\}, \quad \forall i \in I, \forall j \in J. \end{aligned}$$

Equal to 1 if server  $i$  is used

Each service must be deployed on exactly one server

The total amount of resource requested from all services cannot exceed the capacity of the server for each resource



## Round Robin Fashion

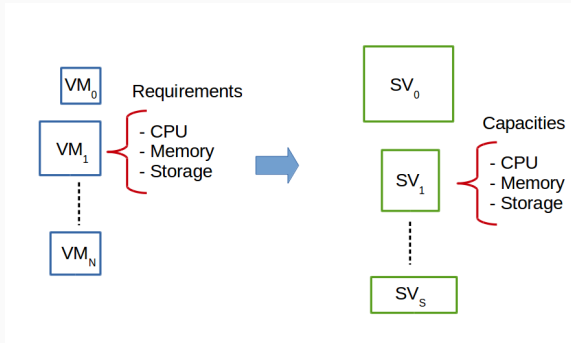
1. Set of requests of VMs deployments
2. Number of servers =  $\lceil \frac{\text{max\_res\_VMs}}{\text{server}_{cap}} \rceil$
3. VMs requests distributed in a Round Robin manner to the appropriate number of servers

## Round Robin Fashion

1. Set of requests of VMs deployments
2. Number of servers =  $\lceil \frac{\text{max\_res\_VMs}}{\text{server}_{cap}} \rceil$
3. VMs requests distributed in a Round Robin manner to the appropriate number of servers

**Drawbacks?**

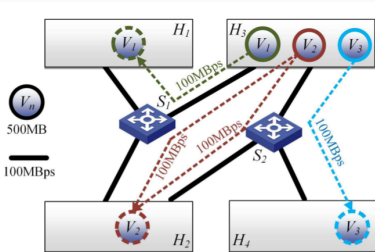
# Multidimensional Case



## Dual Vector Bin-Packing

- Optimization approaches: not scalable
- Optimization phase: expensive in terms of time
- A controller can respond to overload or underload on a server during runtime and reallocates
- Possible through the service **live migration**

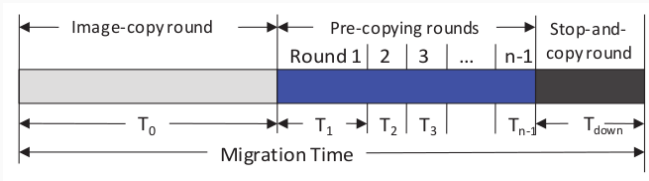
# Live Migration



- “Moving VMs from one physical machine to another without disrupting services”<sup>a</sup>
- High availability
- Minimize the downtime
- Minimize the number of SLA violations
- Virtual machines: stateful

<sup>a</sup>H. Wang, Y. Li, Y. Zhang, D. Jin, “Virtual machine software-defined networks”, IEEE Transactions on Cloud

# Service Migration Phases (VMs)



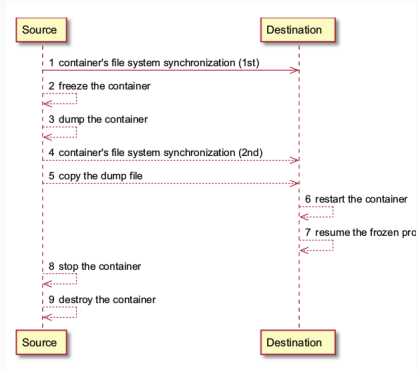
- Data volume transmitted at each round:

$$V_i = \begin{cases} M, & i = 0 \\ R \cdot T_{i-1}, & \text{otherwise} \end{cases}$$

• M: VM image  
• R: copy rate

- Total migration time
- Downtime

# Containers



- A migration is triggered if the utilization of a server exceeds or falls below a certain threshold
- Monitoring system to receive the CPU and memory load of all servers and VMs in a three-second interval (typically)
- Data recorded and stored in a buffer for ten minutes
- Overload and underload situations detected by a control process running every five minutes
- **Thresholds**
- **Migration target servers:** chosen based on their volume



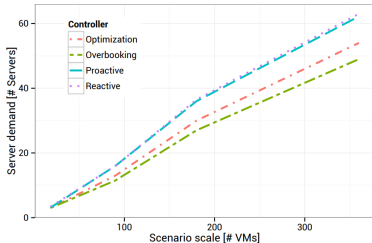
# Proactive Control

- Extension of reactive control
- **Avoid unnecessary migrations**
- **Migration:** only triggered if the forecast suggests that the overload and underload continues and is not only driven by an unforeseen spike in the demand
- For each server a load forecast is computed using one minute of utilization measurements
- If the forecast and  $M$  out of  $K$  measurements pass a threshold an overload or underload situation is detected
- **Forecast:** computed using *double exponential smoothing (DES)*

# Static vs Dynamic

**Dynamic allocation:** often seen as the most efficient means to allocate resources in a data center. Is it always true?

**Experimental results:**



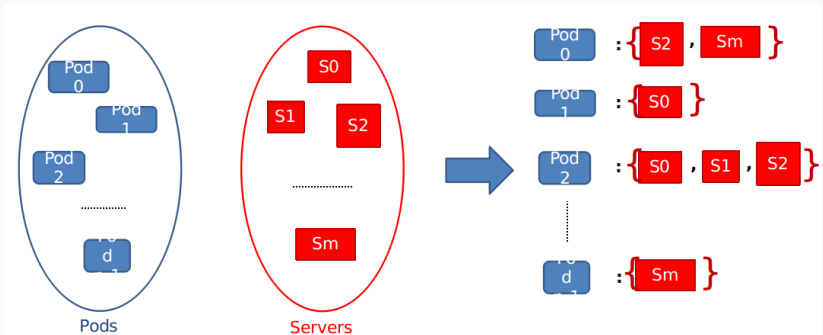
**Suggestions:** good combination of both mechanisms

- **Static allocation:** long term period allocation
- **Dynamic:** exceptional workload peaks

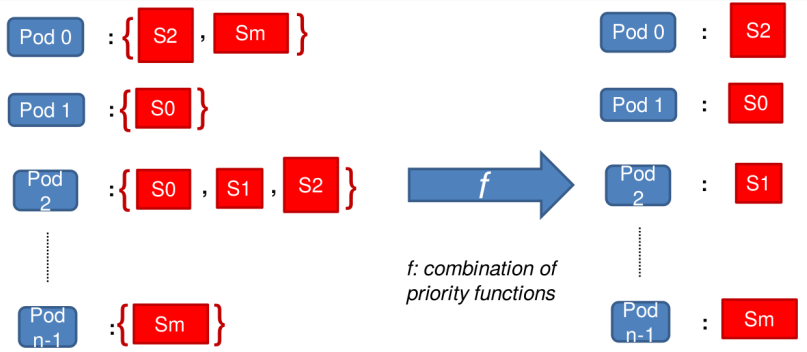
## Well-known cases

- **Kubernetes:** “Open-source system for automatic deployment, scaling, and management of containerized applications”
- **Pod:** “Group of one or more containers with shared storage/network”
- **Heuristic solutions:**
  - **Input:** Set of servers of a cloud infrastructure; set of pods to be deployed;
  - Three main steps:
    1. *Filtering*
    2. *Ranking*
    3. *Deployment*

# Filtering



# Ranking



# Priority Functions

- **LeastRequestPriority**: prefers the node with highest free fraction (in terms of CPU and memory)
- **CalculateNodeLabelPriority**: prefers nodes with specified labels
- **BalancedResourceAllocation**: CPU and Memory balanced after the Pod deployment
- **CalculateSpreadPriority**: minimizes the number of pods belonging to the same service on the same node
- **CalculateAntiAffinityPriority**: minimizes the number of pods of the same service on nodes with the same value for a particular label

# Resource allocation in the Fog





## Problem (1)

# Problem (1)

- Cloud: far  
(high  
latency)



# Problem (1)

- Cloud: far  
(high  
latency)



- High costs



# Problem (1)

- Cloud: far (high latency)



- High costs



- Privacy issues



# Problem (1)

- Cloud: far (high latency)



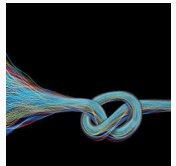
- High costs



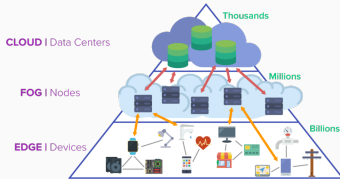
- Privacy issues



- Network congestions

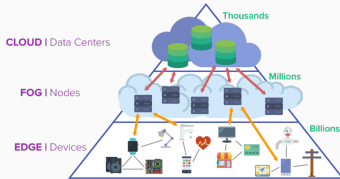


# Fog Computing



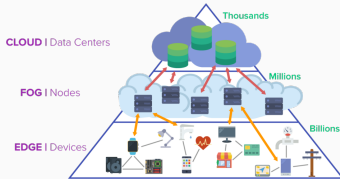
- Extension of cloud computing
- Fog nodes: access points, routers, switches, base stations, servers
- Computation closer to things
- Geographical distribution

# Fog Computing



- Reduction of data traffic towards the cloud
- Extension of cloud computing
- Fog nodes: access points, routers, switches, base stations, servers
- Computation closer to things
- Geographical distribution

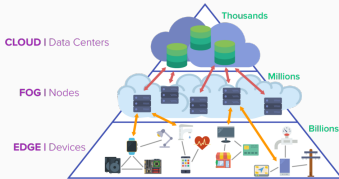
# Fog Computing



- Extension of cloud computing
- Fog nodes: access points, routers, switches, base stations, servers
- Computation closer to things
- Geographical distribution
- Reduction of data traffic towards the cloud
- Low latency

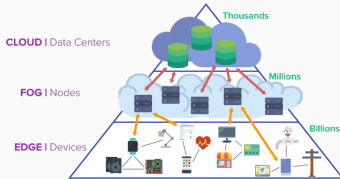


# Fog Computing



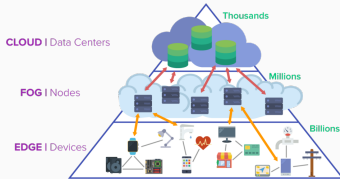
- Extension of cloud computing
- Fog nodes: access points, routers, switches, base stations, servers
- Computation closer to things
- Geographical distribution
- Reduction of data traffic towards the cloud
- Low latency
- Support for mobility

# Fog Computing

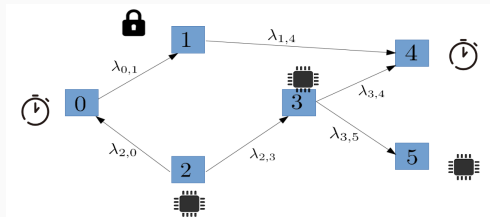


- Extension of cloud computing
- Fog nodes: access points, routers, switches, base stations, servers
- Computation closer to things
- Geographical distribution
- Reduction of data traffic towards the cloud
- Low latency
- Support for mobility
- Reduced costs

# Fog Computing

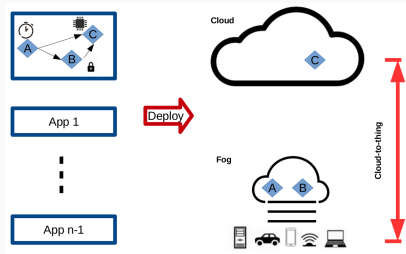


- Extension of cloud computing
- Fog nodes: access points, routers, switches, base stations, servers
- Computation closer to things
- Geographical distribution
- Reduction of data traffic towards the cloud
- Low latency
- Support for mobility
- Reduced costs
- Privacy improved



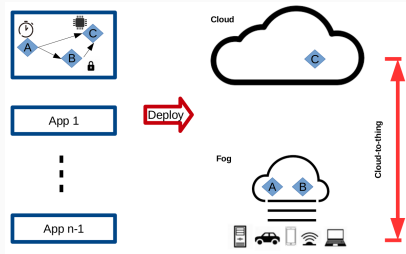
- Microservice architecture
- Multiple coupled modules
- Different requirements (CPU, memory, storage and bandwidth)

# Challenge



- Application modules:  
displaced on the  
cloud-to-things continuum
- Heterogeneous nodes

# Challenge



- Application modules:  
displaced on the  
cloud-to-things continuum
- Heterogeneous nodes



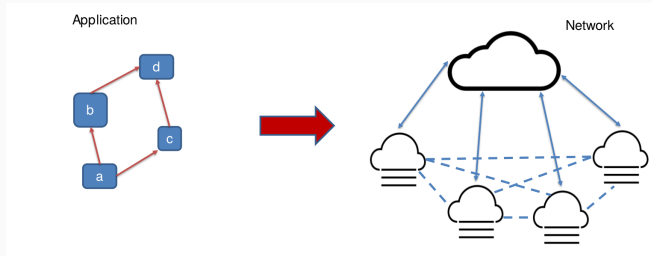
# Resource Allocation in Fog

- Fundamental role in the trade-off between resource utilization and application requirements satisfaction
- **Two approaches:**
  - Static
  - Dynamic
- **Two players**
  - Infrastructure owner
    - *Objective:* maximize the revenue, maximizing the number of applications deployed on the infrastructure
  - Application owner
    - *Objective:* minimize the deployment cost and the completion time

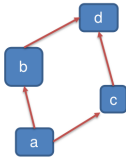
- Applications and Network graphs
- Application deployment reduces to Virtual Network Embedding (VNE) problem (variant of Subgraph Isomorphism): NP-hard
- Approximation algorithms
- Heuristics:
  - Node ranking for each application module
  - Routing between deployed modules



# Static Allocation

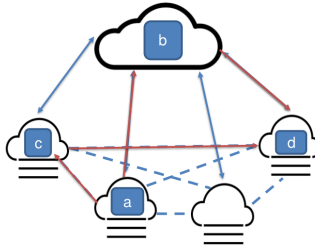


# Static Allocation



Application sorting:

- {a,b,c,d}
- {a,c,b,d}
- {a,b,d,c}
- {a,c,d,b}

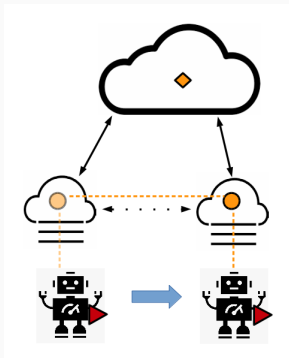


1. Placement (Node Mapping)
2. Routing (Link Mapping)

- Different motivations w.r.t. the Cloud
- Live migration is applied to:
  - Deal with the geographical distribution
  - Save money on the cloud
  - Better utilize fog resources
  - Guarantee a good level of QoS

# Dynamic Allocation

## Industrial scenario

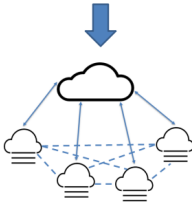
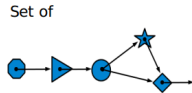


- Monitoring the nodes occupation
- Monitoring clouds costs
- Verifying the resource availability of target nodes
- Perform the migration

# Static vs Dynamic in Fog

- Complete knowledge of the application requirements and network availability
- Optimization problem
- NP-hard problems
- Approximated solutions
- Partial knowledge
- Variability of applications performances
- Live migration
- Monitoring systems
- Application scaling

# Research (Static Allocation)

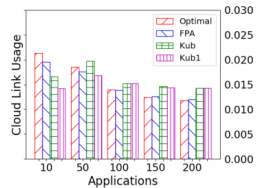
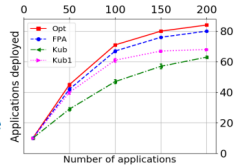


## Constraints:

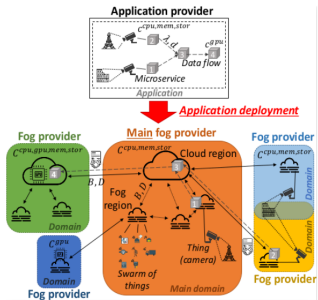
- CPU, RAM, Storage
- Bandwidth
- Delay

**Objective:** maximize profit

**NP-hard problem!**



# Federated Fog

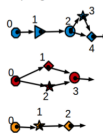


## Constraints:

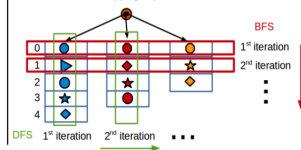
- CPU, RAM, Storage
- Bandwidth
- Delay
- Locality constraints

**Objective:** minimize cost

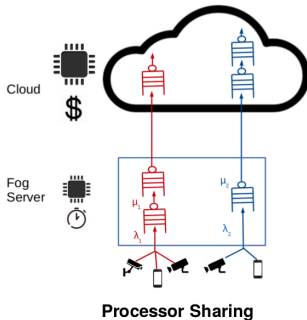
## Topological Sort



## Node Mapping Exploration



# Dynamic Allocation

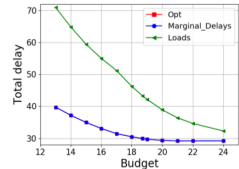
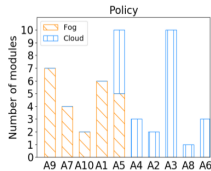


## Constraints:

- Budget
- Stability

**Objective:** minimize total average delay

**Marginal delays:** increase of processing delay in Fog





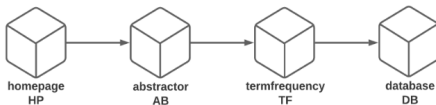
# Hybrid Scenario

Hybrid Cloud environment



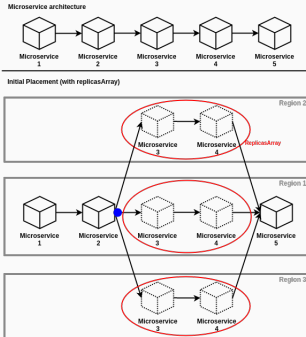
Limited resources at the edge

Microservice-based Application  
(paper miner)

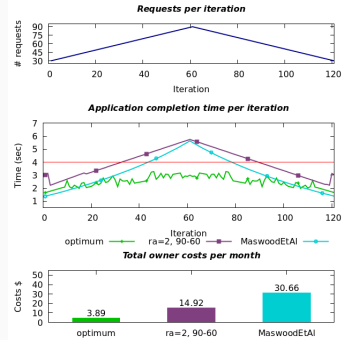


**Objective:** Minimize public cloud cost while guaranteeing application's QoS requirements

# Cost-Effective Workload Allocation Strategy



- Orchestration problem with replicas
- NP-hard problem



- Trade-off: Cost vs. Completion Time

# Take-home Message

- Resource allocation: “old-fashioned” problem
- Strong theoretical foundations
- Fundamental role in distributed systems
- Different problems depending on the context
- Fog and Cloud: new challenges and objectives
- New algorithmic solutions

## Readings:

- A. Wolke, Ma. Bichler, and T. Setzer «Planning vs. Dynamic Control: Resource Allocation in corporate Clouds», IEEE Transactions on Cloud Computing, 2016
- B. Speitkamp and M. Bichler, «A mathematical programming approach for server consolidation problems in virtualized data centers», IEEE Trans. Services Comput., 2010
- S. Akoush, R. Sohan, A. Rice, A.W. Moore, and A. Hopper, «Predicting the performance of virtual machine migration», in Proc. Int. Symp. Model.,Analys.,Simul.,Comput. Syst., 2010

## Additional readings:

- M. Sipser, «Introduction to the Theory of Computation» (Book)
- F. S. Hillier, G. J. Lieberman, «Introduction to Operations Research» (Book)
- T. Ibaraki and N. Katoh, «Resource Allocation Problems: Algorithmic Approaches» (Book)



“Nebulat ergo cogito”