

# Fog and Cloud Computing *Lab*

Daniele Santoro ([dsantoro@fbk.eu](mailto:dsantoro@fbk.eu)) - Expert Research Engineers  
Silvio Cretti ([scretti@fbk.eu](mailto:scretti@fbk.eu)) – Senior Research Engineers

***RiSING (Robust and Secure Distributed Computing)***  
**Fondazione Bruno Kessler (FBK)**

Trento, May 6<sup>th</sup> 2022

# Lab Resources

- Shared Etherpad: <https://annuel2.framapad.org/p/6s5u416vo7-9t4b>
- White Board: <https://tinyurl.com/2p8j7yra>
- Interaction:
  - Etherpad
    - *Exercises check, Share Troubleshooting, Questions and Logs*
  - Zoom Chat (for those remotely connected)
    - *Discuss with your colleagues during exercises or directly/privately with me*
  - Rise your Hand (also via Zoom)
    - *If you need my attention or want to speak, don't be shy !!!*
  - Course Forum: <https://tinyurl.com/27vmd9pj>
    - *Questions and answers could be useful to others, be collaborative*

# Lab Resources


- Slides
  - Uploaded before any lesson in Moodle
- Repositories of exercises
  - <https://gitlab.fbk.eu/dsantoro/fcc-lab-2022>
- Lab Virtual Machine:
  - **Lab VM on Azure (reference for exercises)**
  - Vagrant and VirtualBox on your laptop (possible choice)
    - <https://www.virtualbox.org/>, <https://www.vagrantup.com/> and <https://gitlab.fbk.eu/dsantoro/fcc-lab-2022>

# Quick Recap & Today Lesson

- Recap of previous topics
  - Install a (single-node) k8s cluster in kind
  - Objects in k8s
  - Pod, ReplicaSet, Deployment (missing exercises)
- Multi-Node k8s cluster & Networking
  - Install a multi-node cluster
  - Overlay Network
  - Pod-to-Pod communication
  - External-World-to-Pod communication
  - K8s Services
  - Load Balancer

# From Single-Node to Multi-Node Cluster

## Single Node Cluster




This is also a worker node because kind sets it as a worker by default. Usually, in production clusters the master node does not host any workload.

This is driven by a **taint** on the node:  
`node-role.kubernetes.io/master:NoSchedule`



kubectl



## Multi Node Cluster

## Exercise 24 – Create a multi-node k8s cluster

- **Time:** ~5 minutes
  - 5 minutes: *Altogether, Check, Verify, Ask*

**Description:** Delete the current cluster which is composed by a single worker node (the master node) and create a fresh new cluster composed by one master node and two worker nodes. Check cluster and nodes status.

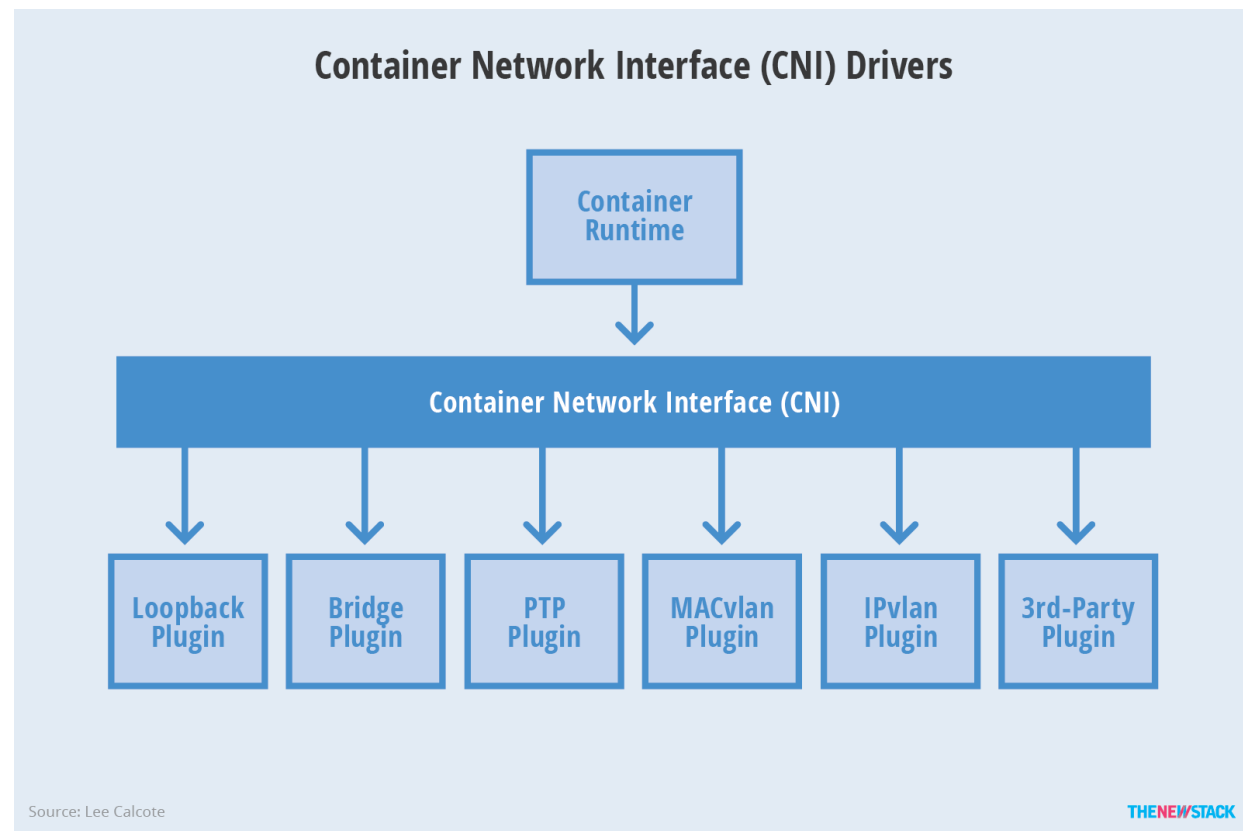
- **Instructions:**  
<https://gitlab.fbk.eu/dsantoro/fcc-lab-2022/-/tree/master/e24>

# Networking Requirements

- To have a fully functional cluster, during installation phase, we need to make sure of the following requirements:
  - A unique IP address is assigned to each Pod
  - Containers in a Pod can communicate to each other
  - The Pod is able to communicate with other Pods in the cluster
  - If configured, the application deployed inside a Pod is accessible from the external world.
- All of the above are networking challenges which must be addressed.

# Container Network Interface (CNI)


- In Kubernetes, each Pod gets a unique IP address.
- For container networking, there are two primary specifications:
  - Container Network Model (CNM), proposed by Docker
  - Container Network Interface (CNI), proposed by CoreOS (used by k8s)
- The Container Runtime offloads the IP assignment to CNI, which connects to the underlying configured plugin, like Bridge or MACvlan, to get the IP address. Once the IP address is given by the respective plugin, CNI forwards it back to the requested Container Runtime.





# Container-to-Container Communication Inside a Pod

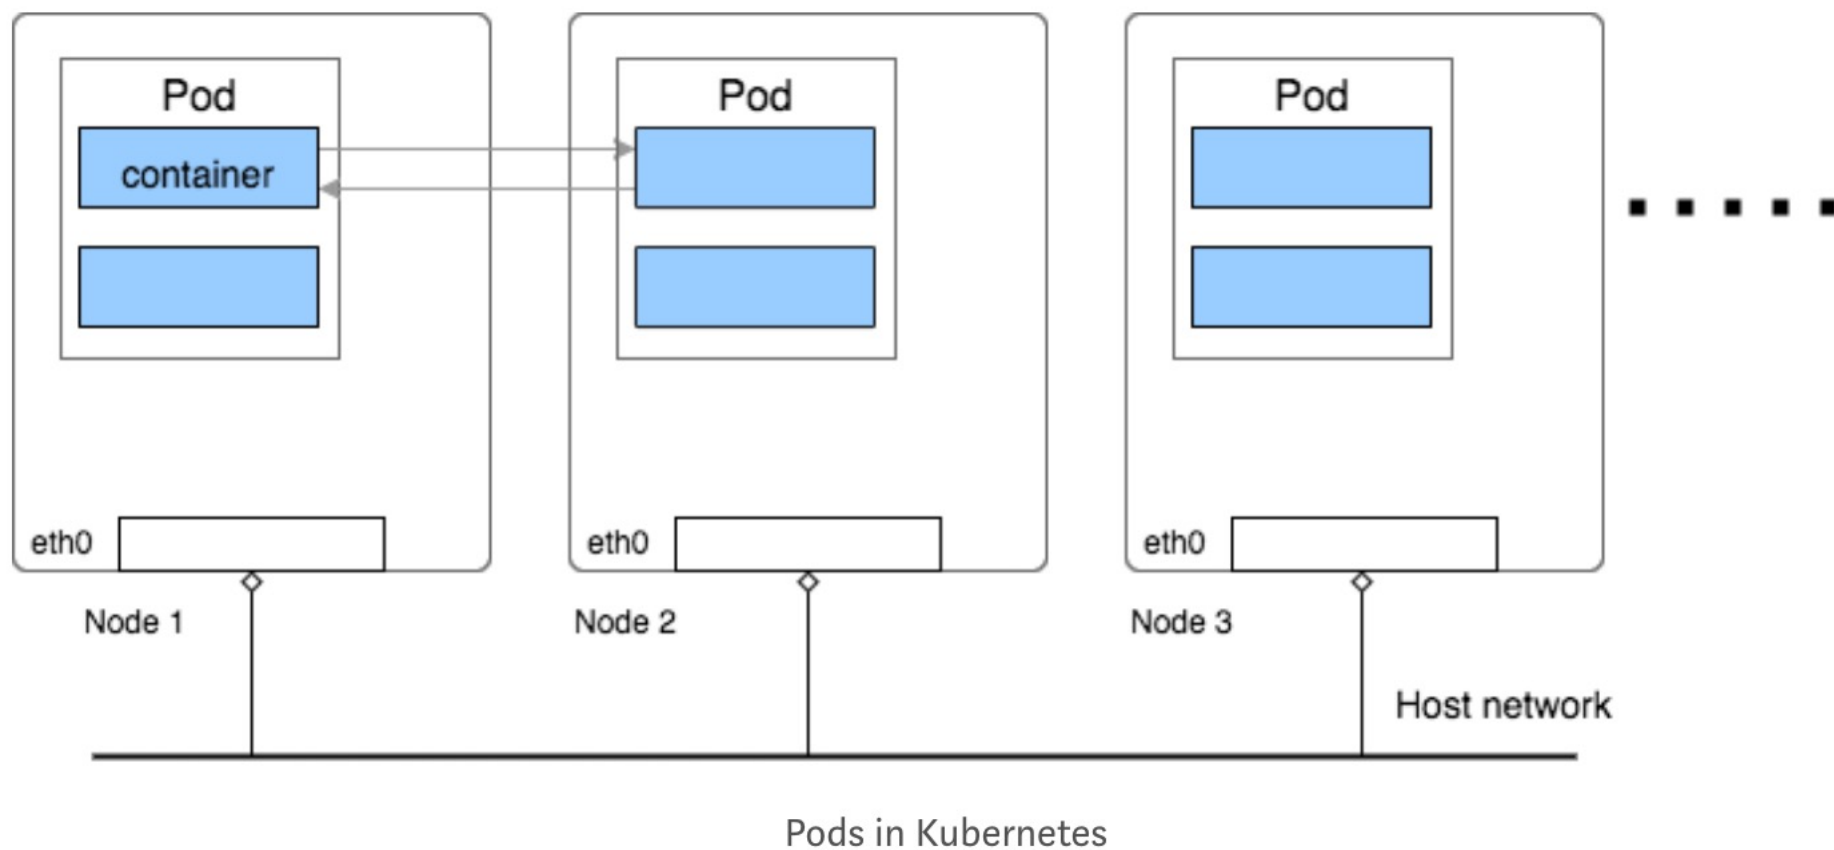
- With the help of the underlying Host OS, all of the Container Runtimes generally create an isolated network entity for each container that they starts.
- On Linux, that entity is referred to as a Network Namespace. These Network Namespaces can be shared across containers, or with the Host Operating System.
- Inside a Pod, containers share the Network Namespaces, so that they can reach to each other via localhost.
- pause is a special container that provides a virtual network interface for the other containers to communicate.



# Pod-to-Pod Communication Across Nodes

- In a clustered environment, the Pods can be scheduled on any node.
- We need to make sure that the Pods can communicate across the nodes, and all the nodes should be able to reach any Pod.
- Kubernetes also puts a condition that there shouldn't be any Network Address Translation (NAT) while doing the Pod-to-Pod communication across Hosts. We can achieve this via:
  - Routable Pods and nodes, using the underlying physical infrastructure, like Google Container Engine
  - Using Software Defined Networking, like [Flannel](#), [Weave](#), [Calico](#), etc.

# k8s Networking



## Exercise 25 - Pod-to-Pod Communications

- **Time:** ~15 minutes
  - 7 minutes: *Try by yourself*
  - 8 minutes: *Check, Verify, Ask*

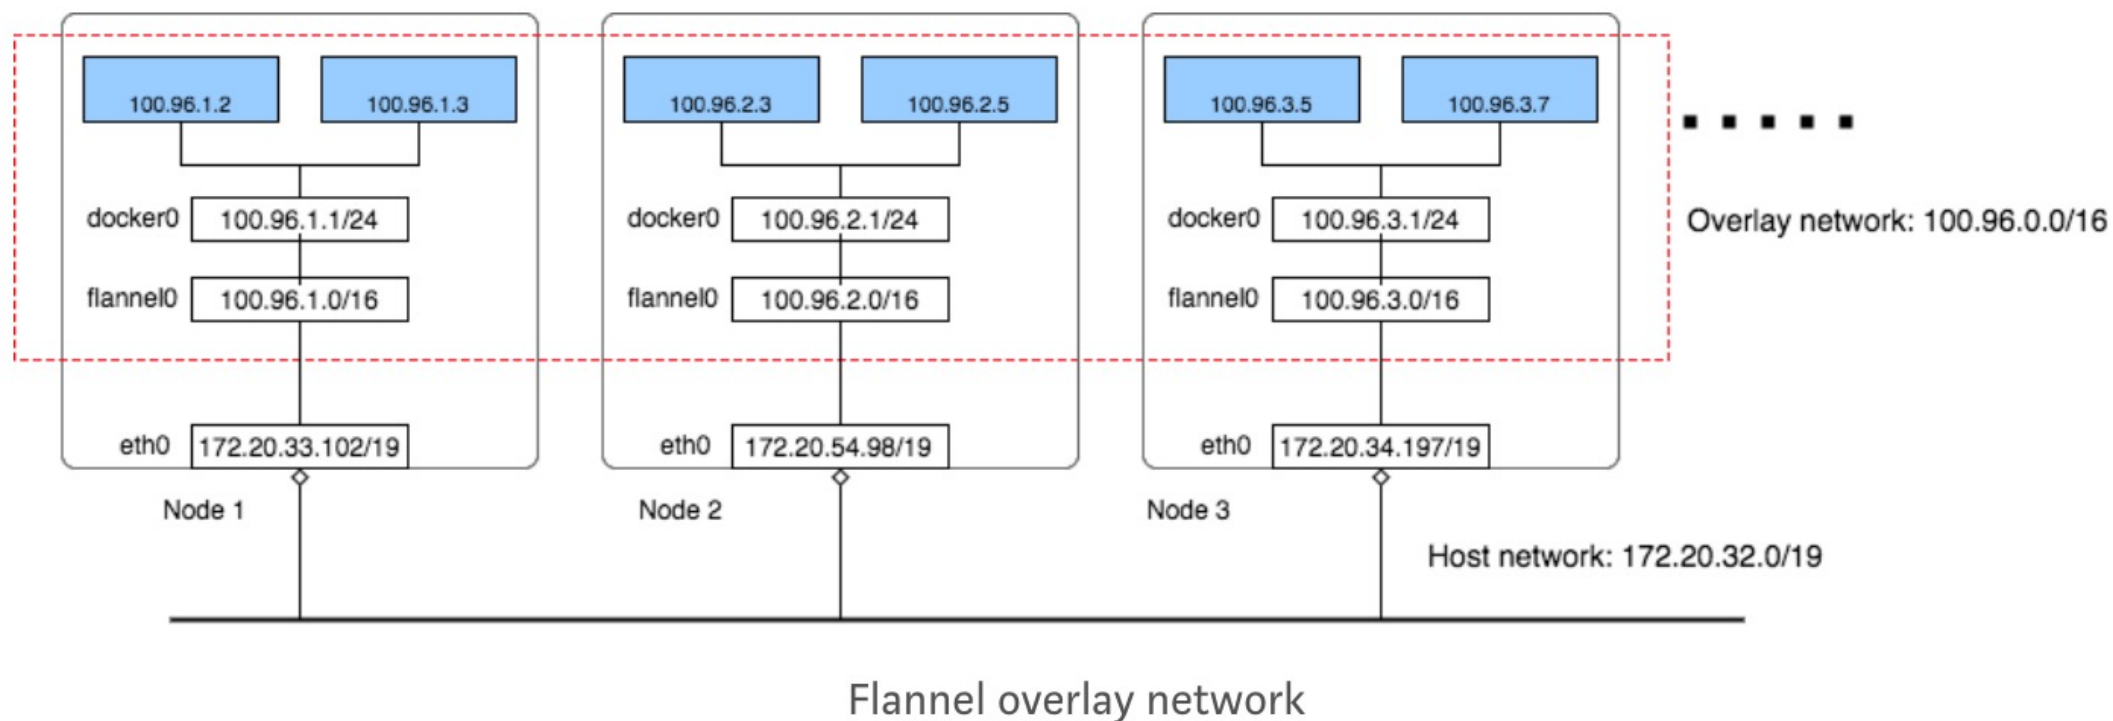
**Description:** Deploy two microservices on different worker nodes: a client and a server. To ensure the scheduler respects your requirements temporarily disable the scheduling on one worker node before installing the second microservice.

Test the communication from client to server ensuring the inter-cluster communications across Pods is working as expected.

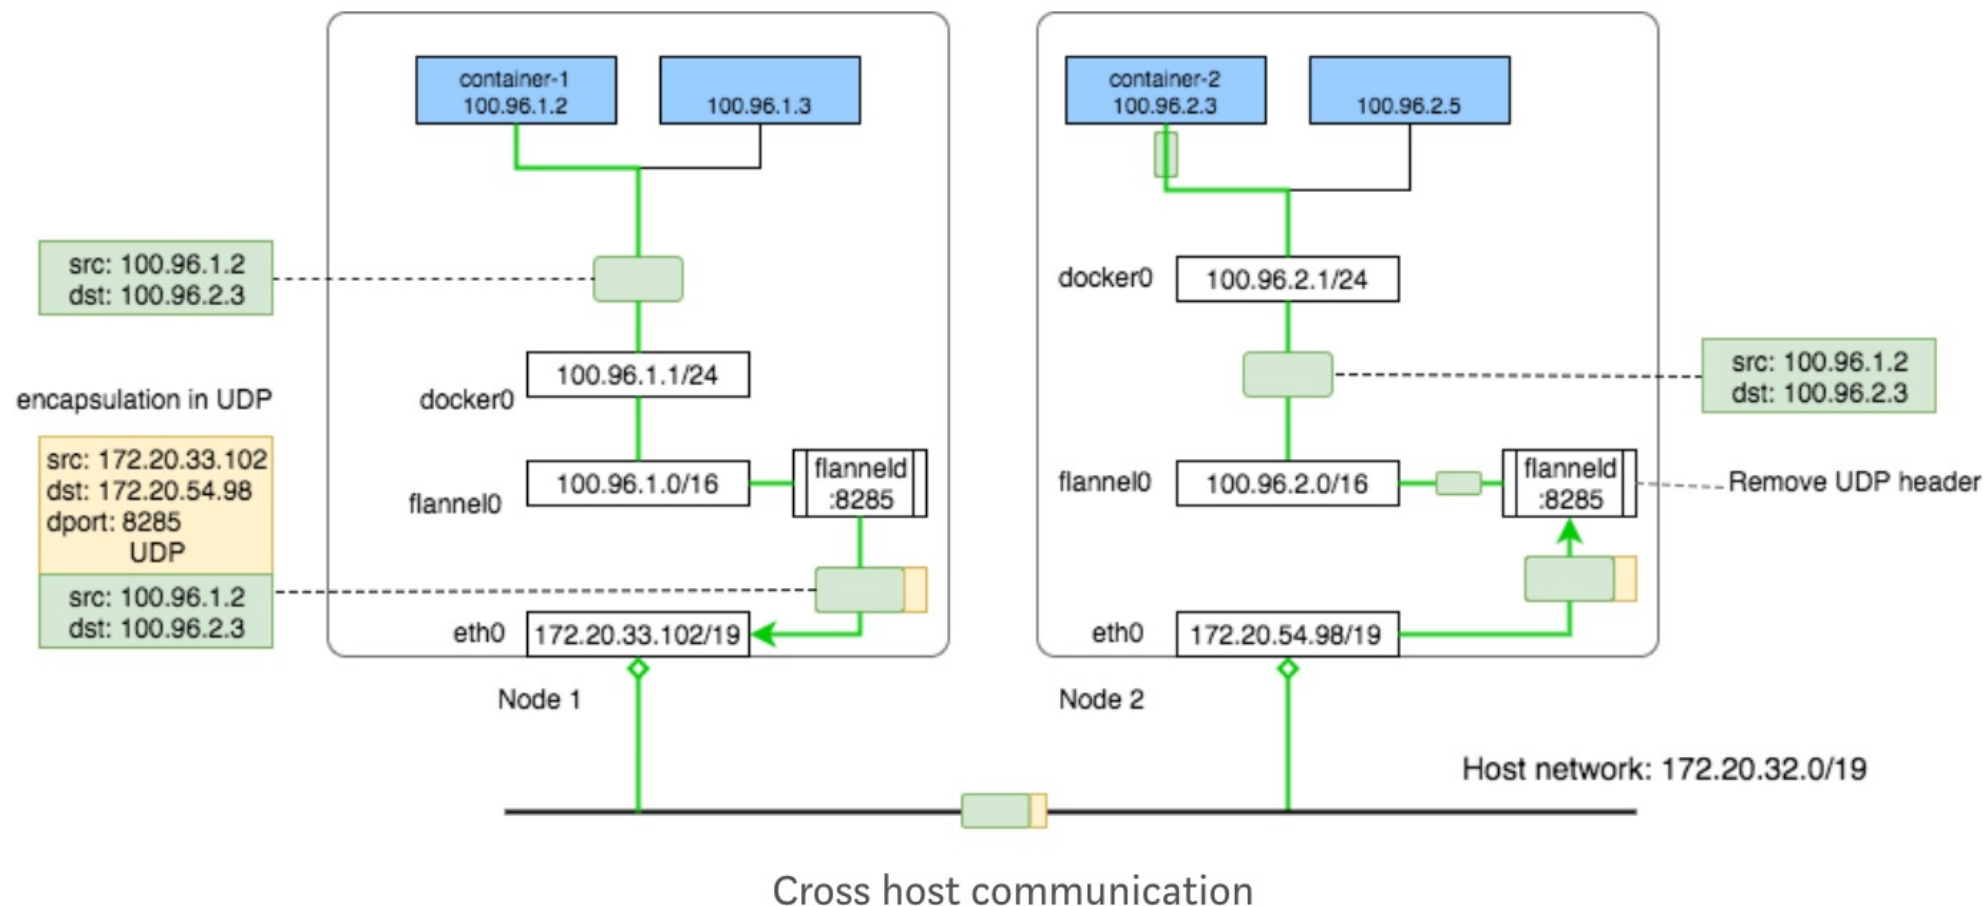
- **Instructions:**

<https://gitlab.fbk.eu/dsantoro/fcc-lab-2022/-/tree/master/e25>

# Overlay Network



# Network Plugin - Flannel




Flannel uses UDP encapsulation in order to encapsulate generic packets into UDP packets.  
Many implementations: IPSec, VXLAN, others

# External World-to-Pod Communication

- By exposing our services to the external world with **kube-proxy**, we can access our applications from outside the cluster.
- Service [[ref](#)] is an high-level abstraction, which logically groups Pods and add policies to access them. This grouping is achieved via Labels and Selectors

# Services


- Scenario in which a user/client is connected to a Pod using its IP address.
- Unexpectedly, the Pod to which the user/client is connected dies, and a new Pod is created by the controller.
- The new Pod will have a new IP address, which will not be known automatically to the user/client of the earlier Pod.

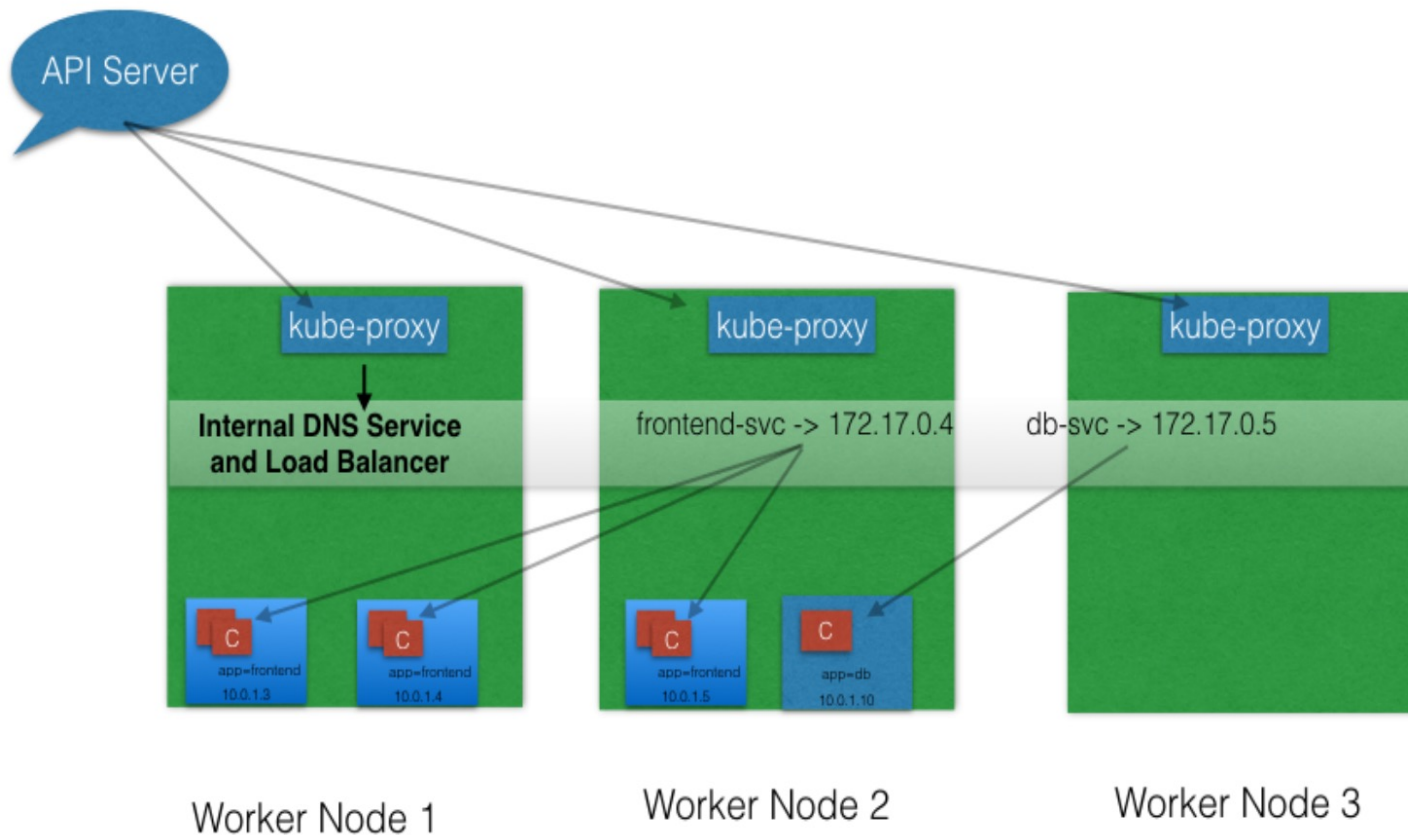




# Services 1/3

- Kubernetes provides a higher-level abstraction called [Service](#), which logically groups Pods and a policy to access them. This grouping is achieved via Labels and Selectors.
- Using Selectors (**app==frontend** and **app==db**), we can group them into two logical groups: one with 3 Pods, and one with just one Pod.
- We can assign a name to the logical grouping, referred to as a **service name**. In our example, we have created two Services, **frontend-svc** and **db-svc**, and they have the **app==frontend** and the **app==db** Selectors, respectively.





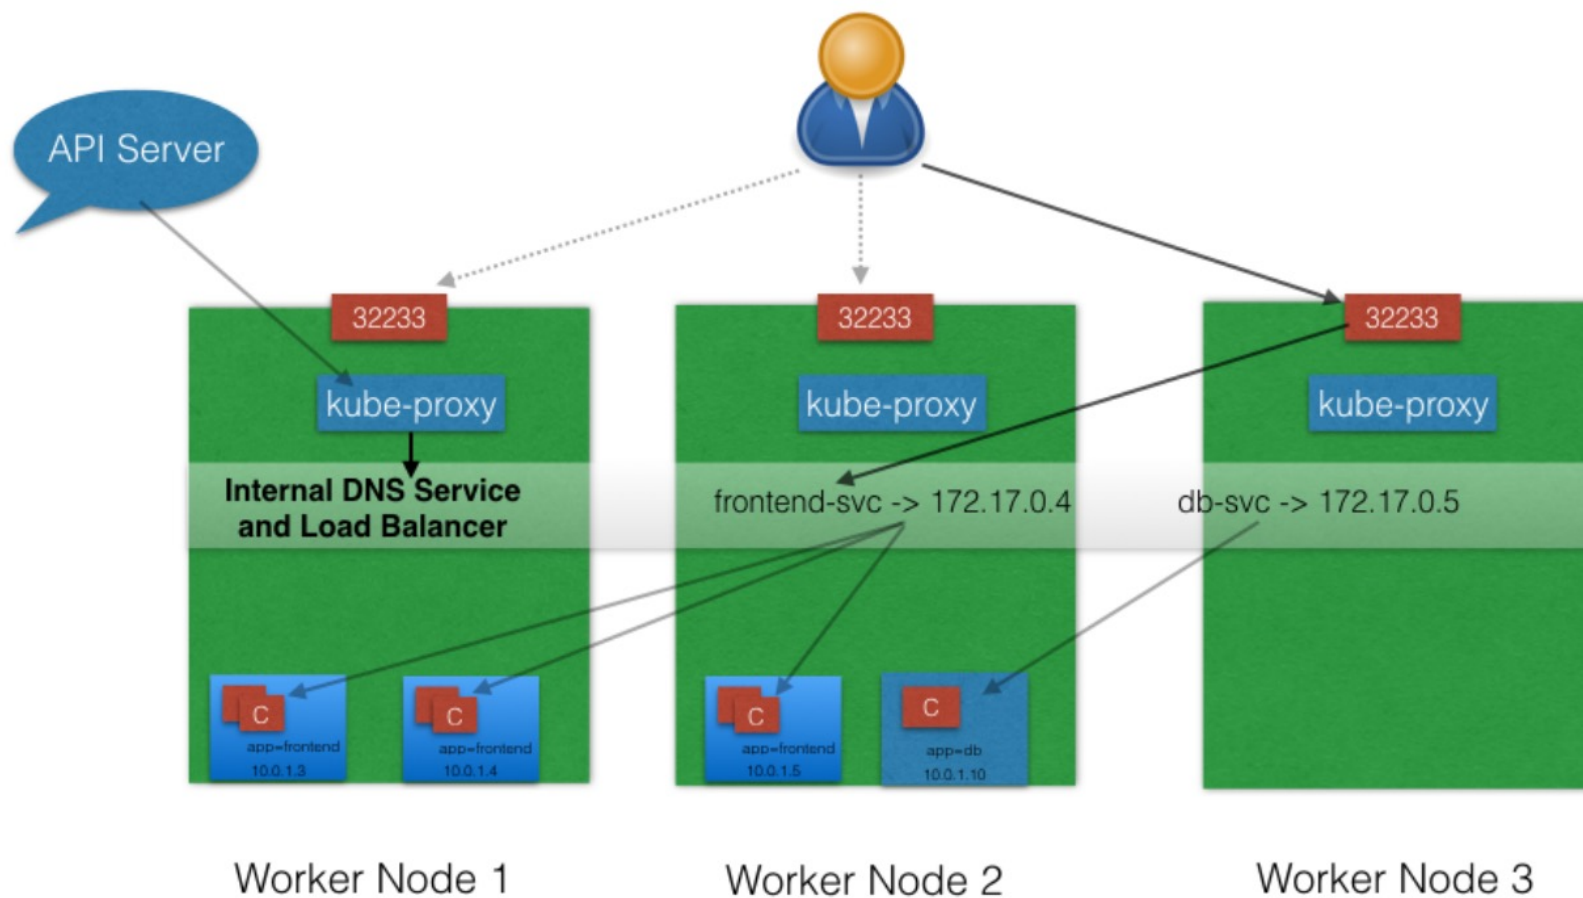
## Services 3/3

- All of the Worker Nodes run a daemon called **kube-proxy**, which watches the API Server on the Master Node for the addition and removal of Services and endpoints.
- For each new Service, on each node, **kube-proxy** configures the IPTables rules to capture the traffic for its ClusterIP and forwards it to one of the endpoints. When the Service is removed, **kube-proxy** removes the IPTables rules on all nodes as well.

# Type of Services

- **ClusterIP:** Exposes the service on a cluster-internal IP. Choosing this value makes the service only reachable from within the cluster. This is the default ServiceType.
- **NodePort:** Exposes the service on each Node's IP at a static port (the NodePort). A ClusterIP service, to which the NodePort service will route, is automatically created. You'll be able to contact the NodePort service, from outside the cluster, by requesting <NodeIP>:<NodePort>.
- **LoadBalancer:** Exposes the service externally using a cloud provider's load balancer. NodePort and ClusterIP services, to which the external load balancer will route, are automatically created.
- **ExternalName:** Maps the service to the contents of the externalName field (e.g. foo.bar.example.com), by returning a CNAME record with its value. No proxying of any kind is set up.

# Node Port Service Type



## Exercise 26 – ExternalWorld-to-Pod Communications

- **Time:** ~15 minutes
  - 6 minutes: *Try by yourself*
  - 9 minutes: *Check, Verify, Ask*

**Description:** Deploy a microservices that act as a server. It must expose a service showing the Pod name and the Worker node where it has been scheduled. Moreover this service should be exposed outside the cluster on a specific port, the 30000. Finally expose the very same microservice using a random external port.

- **Instructions:**  
<https://gitlab.fbk.eu/dsantoro/fcc-lab-2022/-/tree/master/e26>

## Exercise 27 - Load Balancing

- **Time:** ~15 minutes
  - 6 minutes: *Try by yourself*
  - 9 minutes: *Check, Verify, Ask*

**Description:** Scale the number of microservices of the server Deployment (of course is the ReplicaSet which scale) created during previous exercise 6 replicas. While scaling it look at the workload in the cluster and especially paying attention on how the scheduler spread the workload. Then have a look at how the Service resource has been modified by Kubernetes and finally try to access the server many times.

- **Instructions:**

<https://gitlab.fbk.eu/dsantoro/fcc-lab-2022/-/tree/master/e27>